

SI 330: DATA MANIPULATION

---

# LARGE-SCALE DISTRIBUTED COMPUTING

---

# COURSE TRAJECTORY

- ▶ Concepts

  - ▶ Big Data

  - ▶ Distributed Computing

- ▶ Technologies

  - ▶ Hadoop: MapReduce, Hive, Pig, Spark

  - ▶ MRJob, pyspark

  - ▶ local, virtual machine, Advanced Research Computing

---

# BIG DATA

- ▶ The Three Vs:
  - ▶ Volume
    - ▶ not just sampling
  - ▶ Velocity
    - ▶ real-time
  - ▶ Variety
    - ▶ text, audio, images, video

---

# DISTRIBUTED COMPUTING

- ▶ vertical vs. horizontal scaling
- ▶ what's your limit?
- ▶ computing time (weeks vs. days vs. hours vs. minutes)

---

# KEY REQUIREMENTS OF DISTRIBUTED COMPUTING

- ▶ Fault tolerance
- ▶ Recoverability
- ▶ Consistency
- ▶ Scalability

---

# HADOOP

- ▶ “The Apache™ Hadoop® project develops open-source software for reliable, scalable, distributed computing”

# HADOOP

- ▶ HDFS
- ▶ YARN
- ▶ MapReduce
- ▶ Ambari
- ▶ Hive
- ▶ Pig
- ▶ Spark




---

## HADOOP'S ANSWER TO KEY REQUIREMENTS

- ▶ data distribution
- ▶ fixed data block size and replication
- ▶ jobs and tasks
- ▶ high-level programming
- ▶ minimizing network traffic
- ▶ redundant jobs
- ▶ master-worker paradigm



# HADOOP ARCHITECTURAL STACK



---

## HDFS: HADOOP DISTRIBUTED FILE SYSTEM

- ▶ getting the data where it needs to be
  - ▶ (i.e. close to where the computing is going to happen)
- ▶ typically 64MB, 128MB or 256MB blocks
- ▶ access is provided by NameNode and Secondary NameNode (think pilot and co-pilot)

---

## YARN: YET ANOTHER RESOURCE NAVIGATOR

- ▶ decouples job and workload management from cluster and resource management
- ▶ YARN is to computing resources as HDFS is to storage

---

## HDFS: A PRIMER

- ▶ getting your data onto a Hadoop machine isn't enough: you need to make it available to computations
- ▶ How? Put it into HDFS!
- ▶ HDFS looks like a plain old file system (but it's not)
- ▶ `hadoop fs` provides an interface to the filesystem that uses familiar syntax

---

# HADOOP FS

- ▶ `hadoop fs -ls`
  - ▶ lists the files in your HDFS directory
- ▶ `hadoop fs -copyFromLocal <filename>`
  - ▶ copies <filename> from your file system to HDFS
- ▶ `hadoop fs -rm <filename>`
  - ▶ deletes <filename> from your HDFS directory
- ▶ `hadoop fs -cat <filename>`
  - ▶ shows contents of <filename>

---

## ABOUT THE CONCEPT OF 'LOCAL'

- ▶ we are going to be using hadoop in three modes:
  - ▶ local (e.g. MRJob): no real hadoop infrastructure required
  - ▶ local virtual machine (e.g. Hortonworks): single-node hadoop infrastructure
  - ▶ Advanced Research Computing (ARC) Flux Cluster

---

## WHY USE LOCAL?

- ▶ Local
  - ▶ Pros: easy to set up (almost none required), good for teaching concepts about MapReduce
  - ▶ Cons: not really Hadoop, can't go beyond MapReduce

---

## WHY USE LOCAL VIRTUAL MACHINE?

- ▶ Pros:

- ▶ Complete Hadoop ecosystem on your laptop

- ▶ Cons:

- ▶ 11GB required for image, plus more for storage
  - ▶ only provides single-node functionality



---

## WHY USE FLUX (ARC)?


- ▶ Pros:

- ▶ Provides complete Hadoop experience

- ▶ Cons:

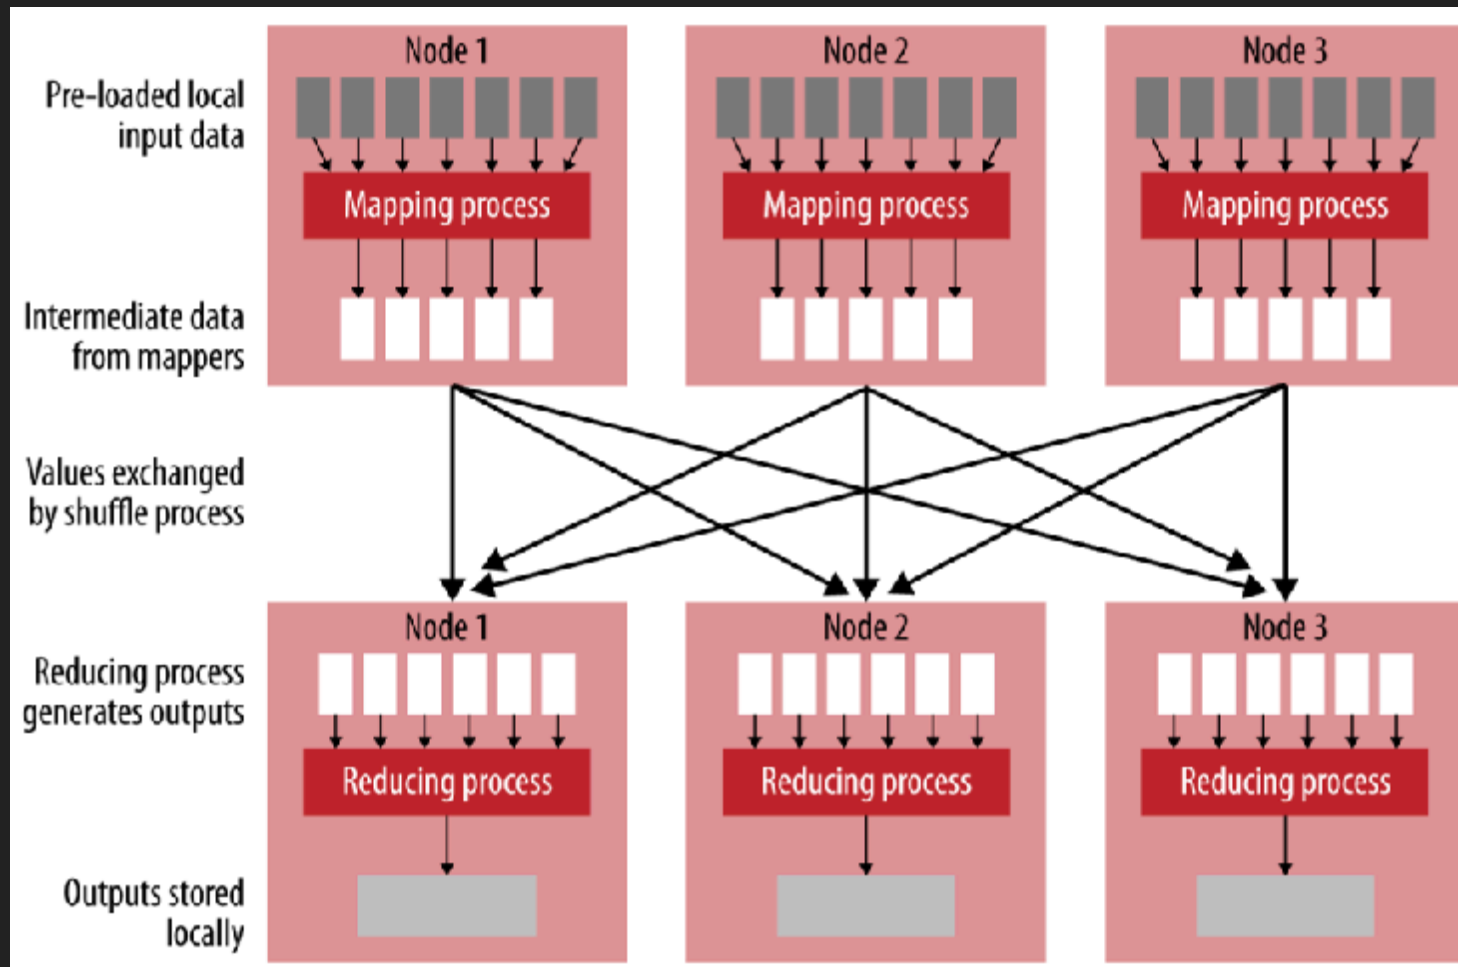
- ▶ Provides complete Hadoop experience

## LET'S LOOK AT MAPREDUCE




- ▶ (book\_id, rating) -> count ratings -> (book\_id, #ratings)

# MAPREDUCE: MAP, SORT & SHUFFLE, REDUCE



# MAPREDUCE: WORD-COUNT EXAMPLE



---

## CLASS QUESTION:

- ▶ What are the necessary conditions to be able to use MapReduce?

---

## TWO WAYS TO TACKLE MAPREDUCE

- ▶ Step-by-step:
  - ▶ You provide the mapper, the reducer, optionally more
  - ▶ mrjob
- ▶ Table-level:
  - ▶ Load, transform, dump (saaaaay....)
  - ▶ High-level SQL-like manipulation of data
  - ▶ Hive, Pig, Spark

---

## LET'S START WITH THE DIY APPROACH: MRJOB

- ▶ `conda install mrjob`
- ▶ stand-alone: can use entirely locally or hook up to more powerful back-end (like Amazon's Elastic MapReduce)  
Text
- ▶ see <http://pythonhosted.org/mrjob> for documentation
- ▶ must define at least one of: mapper, reducer, combiner

# MRJOB FRAMEWORK AND WORD COUNTS

- ▶ Mapper:

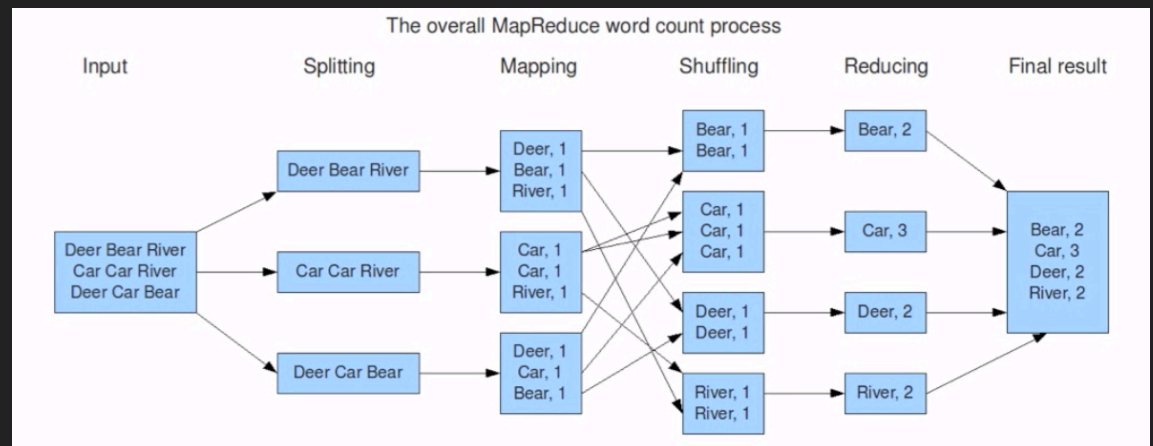
- ▶ break input line into key, value pairs

- ▶ Reducer:

- ▶ take all (key, value) pairs with same key and compute aggregate function (e.g. sum, count)

- ▶ MRJob takes care of:

Sorting mapper output, invoking reduce tasks, assembling final output, scheduling and monitoring tasks





# A BASIC MRJOB MAPREDUCE EXAMPLE

```
#!/usr/bin/python
```

```
from mrjob.job import MRJob
import re
```

```
class MRWordFrequencyCount(MRJob):
```


```
    """ input: self, in_key, in_value """
```

```
    def mapper(self, _, line):
        yield "chars", len(line)
        yield "words", len(line.split())
        yield "lines", 1
```

```
    """ input: self, in_key from mapper, in_value from mapper """
```

```
    def reducer(self, key, values):
        yield key, sum(values)
```

```
if __name__ == '__main__':
    MRWordFrequencyCount.run()
```



# WHOA, WHOA, WHOA!

```
#!/usr/bin/python

from mrjob.job import MRJob
import re

class MRWordFrequencyCount(MRJob):

    ### input: self, in_key, in_value
    def mapper(self, _, line):
        yield "chars", len(line)
        yield "words", len(line.split())
        yield "lines", 1

    ### input: self, in_key from mapper, in_value from mapper
    def reducer(self, key, values):
        yield key, sum(values)

if __name__ == '__main__':
    MRWordFrequencyCount.run()
```

- ▶ yield?
- ▶ it's a generator thing
- ▶ what?

# GENERATORS

```
def f123():  
    yield 1  
    yield 2  
    yield 3  
for item in f123():  
    print item
```

- ▶ A GENERATOR is a lazy iterable object: lazy objects wait to produce results only when they absolutely have to.
- ▶ When f123() is called, it does not return any of the values in the yield statements.
- ▶ It returns a generator object. Also, the function does not really exit - it goes into a suspended state.
- ▶ When the for loop tries to loop over the generator object, the function:
  - ▶ Resumes from its suspended state
  - ▶ Runs until the next yield statement
  - ▶ Returns that yield result as the next item.
- ▶ This happens until the function exits, at which point the generator raises StopIteration, and the loop exits.

---

## MAPPERS AND REDUCERS

- ▶ mappers and reducers are generators supplied by YOU!
- ▶ too weird?
  - ▶ basically, yield means return and then pick up here next time through

# WORDCOUNT BY MAPPING AND REDUCING

```
#!/usr/bin/python
```

```
from mrjob.job import MRJob
import re
```

```
class MRWordFrequencyCount(MRJob):
```

```
    """ input: self, in_key, in_value """
```

```
    def mapper(self, _, line):
```

```
        yield "chars", len(line)
```

```
        yield "words", len(line.split())
```

```
        yield "lines", 1
```

```
    """ input: self, in_key from mapper, in_value from mapper """
```

```
    def reducer(self, key, values):
```

```
        yield key, sum(values)
```

```
if __name__ == '__main__':
```

```
    MRWordFrequencyCount.run()
```

▶ "To be or not to be,  
That is the question."

▶ ("chars", 18), ("words", 6), ("lines", 1)  
("chars", 20), ("words", 4), ("lines", 1)

▶ "chars": [18,20], "words": [6,4], "lines":  
[1,1]

▶ ("chars", 38), ("words", 10), ("lines", 2)

---

# LET'S RUN THAT MRJOB WORDCOUNT PROGRAM

```
↳ python wordcountmr.py andalusian.txt
No configs found; falling back on auto-configuration
No configs specified for inline runner
Running step 1 of 1...
Creating temp directory /var/folders/1r/ht9dh7xn5j1ddwl8p7frd1282yw4hz/T/wordcountmr.cteplovs.20171105.212554.677036
Streaming final output from /var/folders/1r/ht9dh7xn5j1ddwl8p7frd1282yw4hz/T/wordcountmr.cteplovs.20171105.212554.677036/output...
"chars" 655
"lines" 15
"words" 102
Removing temp directory /var/folders/1r/ht9dh7xn5j1ddwl8p7frd1282yw4hz/T/wordcountmr.cteplovs.20171105.212554.677036...
```

---

## WHAT JUST HAPPENED?

- ▶ we ran our wordcount MapReduce python script on our local machine